

```
In [2]: import pandas as pd
import numpy as np
import datetime as dt
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
from yellowbrick.cluster import KElbowVisualizer
from numpy import math
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
```

1. Loading & Reading the Dataset

```
In [3]: df = pd.read_csv('C:\\Users\\PC\\Desktop\\DA PROJECT\\RFMT ANALYSIS\\online_retail_II_da
print('-'*50)
print('Data imported successfully!!!')
df.head(5).style.set_properties(**{"background-color": "#cd5c5c", "color": "black", "bord
```

Data imported successfully!!!

```
Out[3]:
```

	Invoice	StockCode	Description	Quantity	InvoiceDate	Price	Customer ID	Country
0	489434	85048	15CM CHRISTMAS GLASS BALL 20 LIGHTS	12	2009-12-01 07:45:00	6.950000	13085.000000	United Kingdom
1	489434	79323P	PINK CHERRY LIGHTS	12	2009-12-01 07:45:00	6.750000	13085.000000	United Kingdom
2	489434	79323W	WHITE CHERRY LIGHTS	12	2009-12-01 07:45:00	6.750000	13085.000000	United Kingdom
3	489434	22041	RECORD FRAME 7" SINGLE SIZE	48	2009-12-01 07:45:00	2.100000	13085.000000	United Kingdom
4	489434	21232	STRAWBERRY CERAMIC TRINKET BOX	24	2009-12-01 07:45:00	1.250000	13085.000000	United Kingdom

2. Exploring information of the dataset

```
In [4]: pd.set_option('display.max_columns', None)
def data_overview(df, head=5):
    print('SHAPE OF DATASET'.center(125, '-'))
    print('Rows:{}'.format(df.shape[0]))
    print('Columns:{}'.format(df.shape[1]))
    print('HEAD OF DATASET'.center(125, '-'))
    print(df.head(head))
    print('DATA TYPE'.center(125, '-'))
    print(df.dtypes)
    print('MISSING VALUES'.center(125, '-'))
    print(df.isnull().sum()[df.isnull().sum()>0].sort_values(ascending = False))
    print('DUPLICATED VALUES'.center(125, '-'))
    print(df.duplicated().sum())
    print('STATISTICS OF DATA'.center(125, '-'))
    print(df.describe(include="all"))
    print("DATA INFO".center(125, '-'))
```

```
print(df.info())
data_overview(df)
```

-----SHAPE OF DATASET-----

Rows:1067371
Columns:8

-----HEAD OF DATASET-----

	Invoice	StockCode	Description	Quantity	\
0	489434	85048	15CM CHRISTMAS GLASS BALL 20 LIGHTS	12	
1	489434	79323P	PINK CHERRY LIGHTS	12	
2	489434	79323W	WHITE CHERRY LIGHTS	12	
3	489434	22041	RECORD FRAME 7" SINGLE SIZE	48	
4	489434	21232	STRAWBERRY CERAMIC TRINKET BOX	24	

	InvoiceDate	Price	Customer ID	Country
0	2009-12-01 07:45:00	6.95	13085.0	United Kingdom
1	2009-12-01 07:45:00	6.75	13085.0	United Kingdom
2	2009-12-01 07:45:00	6.75	13085.0	United Kingdom
3	2009-12-01 07:45:00	2.10	13085.0	United Kingdom
4	2009-12-01 07:45:00	1.25	13085.0	United Kingdom

-----DATA TYPE-----

Invoice object
StockCode object
Description object
Quantity int64
InvoiceDate object
Price float64
Customer ID float64
Country object
dtype: object

-----MISSING VALUES-----

Customer ID 243007
Description 4382
dtype: int64

-----DUPLICATED VALUES-----

34335

-----STATISTICS OF DATA-----

	Invoice	StockCode	Description	Quantity	\
count	1067371	1067371	1062989	1.067371e+06	
unique	53628	5305	5698	NaN	
top	537434	85123A	WHITE HANGING HEART T-LIGHT HOLDER	NaN	
freq	1350	5829	5918	NaN	
mean	NaN	NaN	NaN	9.938898e+00	
std	NaN	NaN	NaN	1.727058e+02	
min	NaN	NaN	NaN	-8.099500e+04	
25%	NaN	NaN	NaN	1.000000e+00	
50%	NaN	NaN	NaN	3.000000e+00	
75%	NaN	NaN	NaN	1.000000e+01	
max	NaN	NaN	NaN	8.099500e+04	

	InvoiceDate	Price	Customer ID	Country
count	1067371	1.067371e+06	824364.000000	1067371
unique	47635	NaN	NaN	43
top	2010-12-06 16:57:00	NaN	NaN	United Kingdom
freq	1350	NaN	NaN	981330
mean	NaN	4.649388e+00	15324.638504	NaN
std	NaN	1.235531e+02	1697.464450	NaN
min	NaN	-5.359436e+04	12346.000000	NaN
25%	NaN	1.250000e+00	13975.000000	NaN
50%	NaN	2.100000e+00	15255.000000	NaN

```

75%      NaN      4.150000e+00      16797.000000      NaN
max      NaN      3.897000e+04      18287.000000      NaN
-----DATA INFO-----
-----
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1067371 entries, 0 to 1067370
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Invoice          1067371 non-null  object
1   StockCode       1067371 non-null  object
2   Description     1062989 non-null  object
3   Quantity        1067371 non-null  int64
4   InvoiceDate      1067371 non-null  object
5   Price           1067371 non-null  float64
6   Customer ID     824364 non-null   float64
7   Country         1067371 non-null  object
dtypes: float64(2), int64(1), object(5)
memory usage: 65.1+ MB
None

```

NOTES:

- The Dataset has Rows: 1067371 and Columns:8
- The Dataset has 3 types of columns: strings(5), integer(1), float(2)
- The Dataset has Missing values in Customer ID (243007) and Description (4382)
- Invoice starts with the 'c' needs to be cleaned as it is cancelled transaction
- The Dataset has duplicates
- Also check for negative value and outliers in Quantity and Price

3. Data Wrangling

```

In [5]: df = df.rename(columns = {
        'Customer ID' : 'CustomerID'
    })

```

```

In [6]: # 3.1 Checking the data types:
df['InvoiceDate'] = pd.to_datetime(df['InvoiceDate'])
df.dtypes

```

```

Out[6]: Invoice          object
StockCode       object
Description     object
Quantity        int64
InvoiceDate     datetime64[ns]
Price           float64
CustomerID      float64
Country         object
dtype: object

```

```

In [7]: # 3.2 Dealing with missing values
print("Shape of data before removing NaN's CustomerID",df.shape)
df.dropna(subset="CustomerID",axis=0,inplace=True)
print("Shape of data after removing NaN's CustomerID",df.shape)

```

```

Shape of data before removing NaN's CustomerID (1067371, 8)
Shape of data after removing NaN's CustomerID (824364, 8)

```

```

In [8]: print("Missing values in each column after cleaning customerID :\n", df.isnull().sum())

Missing values in each column after cleaning customerID :

```

```
Invoice      0
StockCode    0
Description  0
Quantity     0
InvoiceDate  0
Price        0
CustomerID   0
Country      0
dtype: int64
```

```
In [9]: # 3.3 Removing canceled products from invoice
df=df[~df['Invoice'].str.contains('C', na=False)]
df.shape
print("Dataset is free from cancelled products information")
```

Dataset is free from cancelled products information

```
In [10]: # 3.4 Removing the duplicates
print('Number of duplicates before cleaning:', df.duplicated().sum())
df=df.drop_duplicates(keep='first')
print('Number of duplicates before cleaning:', df.duplicated().sum())
```

Number of duplicates before cleaning: 26125
Number of duplicates before cleaning: 0

```
In [11]: # 3.5 Checking for negative values
print('Negative values in Quantity is:', (df['Quantity']<0).sum())
print('Negative values in Price is:', (df['Price']<0).sum())
```

Negative values in Quantity is: 0
Negative values in Price is: 0


```
In [12]: # 3.6 Cleaning outliers
def outliers_thresholds(dataframe, variable):
    quartile_1 = dataframe[variable].quantile(0.25)
    quartile_3 = dataframe[variable].quantile(0.75)
    interquartile_range = quartile_3 - quartile_1
    up_limit = quartile_3 + 1.5*interquartile_range
    low_limit = quartile_1 - 1.5*interquartile_range
    return up_limit, low_limit

def replace_with_thresholds(dataframe, variable):
    up_limit, low_limit = outliers_thresholds(dataframe, variable)
    dataframe.loc[(dataframe[variable]<low_limit), variable] = low_limit
    dataframe.loc[(dataframe[variable]>up_limit), variable] = up_limit

print(outliers_thresholds(df, 'Quantity'))
print(outliers_thresholds(df, 'Price'))
```


(27.0, -13.0)
(7.5, -2.5)

```
In [13]: # Observing them before removing outliers.
fig, ax = plt.subplots(2,1, figsize = (20,5))
col_list = ["Quantity", "Price"]
for i in range(0,2):
    ax[i].boxplot(df[col_list[i]], flierprops = dict(marker = "s", markerfacecolor = "red"))
plt.show()
```



```
In [14]: # Apply the function to remove the outliers
replace_with_thresholds(df, 'Quantity')
replace_with_thresholds(df, 'Price')
```

```
In [15]: # Observing them after removing outliers.
fig, ax = plt.subplots(2,1, figsize = (20,5))
for i in range(0,2):
    ax[i].boxplot(df[col_list[i]],flierprops = dict(marker = "s", markerfacecolor = "red"))
plt.show()
```



Data is clean now

4. EDA: Feature Engineer

```
In [16]: #Creating new feature Revenue
df['Revenue'] = df['Quantity']*df['Price']
df
```

	Invoice	StockCode	Description	Quantity	InvoiceDate	Price	CustomerID	Country	Revenue
0	489434	85048	15CM CHRISTMAS GLASS BALL 20 LIGHTS	12	2009-12-01 07:45:00	6.95	13085.0	United Kingdom	83.40
1	489434	79323P	PINK CHERRY LIGHTS	12	2009-12-01 07:45:00	6.75	13085.0	United Kingdom	81.00
2	489434	79323W	WHITE CHERRY LIGHTS	12	2009-12-01 07:45:00	6.75	13085.0	United Kingdom	81.00
3	489434	22041	RECORD FRAME 7" SINGLE SIZE	27	2009-12-01 07:45:00	2.10	13085.0	United Kingdom	56.70
4	489434	21232	STRAWBERRY CERAMIC TRINKET BOX	24	2009-12-01 07:45:00	1.25	13085.0	United Kingdom	30.00

...
1067366	581587	22899	CHILDREN'S APRON DOLLY GIRL	6	2011-12-09 12:50:00	2.10	12680.0	France	12.60
1067367	581587	23254	CHILDRENS CUTLERY DOLLY GIRL	4	2011-12-09 12:50:00	4.15	12680.0	France	16.60
1067368	581587	23255	CHILDRENS CUTLERY CIRCUS PARADE	4	2011-12-09 12:50:00	4.15	12680.0	France	16.60
1067369	581587	22138	BAKING SET 9 PIECE RETROSPOT	3	2011-12-09 12:50:00	4.95	12680.0	France	14.85
1067370	581587	POST	POSTAGE	1	2011-12-09 12:50:00	7.50	12680.0	France	7.50

779495 rows × 9 columns

```
In [17]: print('Max Date:', df['InvoiceDate'].max())
print('Min Date:', df['InvoiceDate'].min())
```

```
Max Date: 2011-12-09 12:50:00
Min Date: 2009-12-01 07:45:00
```

```
In [18]: # Set latest date is 2011-12-10 as the last invoice is on 2011-12-09
Latest_Date = dt.datetime(2011,12,10)
```

```
In [19]: #Creating the RFM features with subsets of CustomerID
RFM = df.groupby(df['CustomerID']).agg({
    'InvoiceDate': lambda x: (Latest_Date - x.max()).days,
    'Invoice': lambda x: x.nunique(),
    'Revenue': lambda x: x.sum()})
RFM.dtypes
```

```
Out[19]: InvoiceDate      int64
Invoice      int64
Revenue      float64
dtype: object
```

```
In [20]: #Renaming column names to Recency, Frequency and Monetary
RFM.rename(columns = {
    'InvoiceDate' : 'Recency',
    'Invoice' : 'Frequency',
    'Revenue' : 'Monetary'}, inplace=True)

RFM.reset_index().head().style.set_properties(**{"background-color": "#cd5c5c","color":
```

```
Out[20]:
```

	CustomerID	Recency	Frequency	Monetary
0	12346.000000	325	12	400.940000
1	12347.000000	2	8	4473.220000
2	12348.000000	75	5	779.730000
3	12349.000000	18	4	3347.990000
4	12350.000000	310	1	301.900000

- The Fourth variable of RFM, InterPurchase Time, is a measure of average time gap between total

shopping trips by a customer.

- The Interpurchase Time is calculated as follows :
- $T = L/F = (T_n - T_1)/F$
- Note: We consider only those customers who made purchase more than once.

```
In [21]: RFM = RFM[RFM['Frequency']>1]
RFM.reset_index().head().style.set_properties(**{"background-color": "#cd5c5c","color":
```

```
Out[21]:
```

	CustomerID	Recency	Frequency	Monetary
--	------------	---------	-----------	----------

0	12346.000000	325	12	400.940000
1	12347.000000	2	8	4473.220000
2	12348.000000	75	5	779.730000
3	12349.000000	18	4	3347.990000
4	12352.000000	36	10	1739.490000

```
In [22]: Shopping_Cycle = df.groupby(df['CustomerID']).agg({'InvoiceDate': lambda x: (x.max() - x
```

```
In [23]: RFM['Shopping_Cycle'] = Shopping_Cycle
RFM.reset_index().head().style.set_properties(**{"background-color": "#cd5c5c","color":
```

```
Out[23]:
```

	CustomerID	Recency	Frequency	Monetary	Shopping_Cycle
--	------------	---------	-----------	----------	----------------

0	12346.000000	325	12	400.940000	400
1	12347.000000	2	8	4473.220000	402
2	12348.000000	75	5	779.730000	362
3	12349.000000	18	4	3347.990000	570
4	12352.000000	36	10	1739.490000	356

```
In [24]: RFM['InterPurchase_Time'] = RFM['Shopping_Cycle'] // RFM['Frequency']
RFM.reset_index().head().style.set_properties(**{"background-color": "#cd5c5c","color":

RFMT = RFM[["Recency", "Frequency", "Monetary", "InterPurchase_Time"]]
RFMT.reset_index().head().style.set_properties(**{"background-color": "#cd5c5c","color":
```

```
Out[24]:
```

	CustomerID	Recency	Frequency	Monetary	InterPurchase_Time
--	------------	---------	-----------	----------	--------------------

0	12346.000000	325	12	400.940000	33
1	12347.000000	2	8	4473.220000	50
2	12348.000000	75	5	779.730000	72
3	12349.000000	18	4	3347.990000	142
4	12352.000000	36	10	1739.490000	35

RFMT Model is ready for segmentation

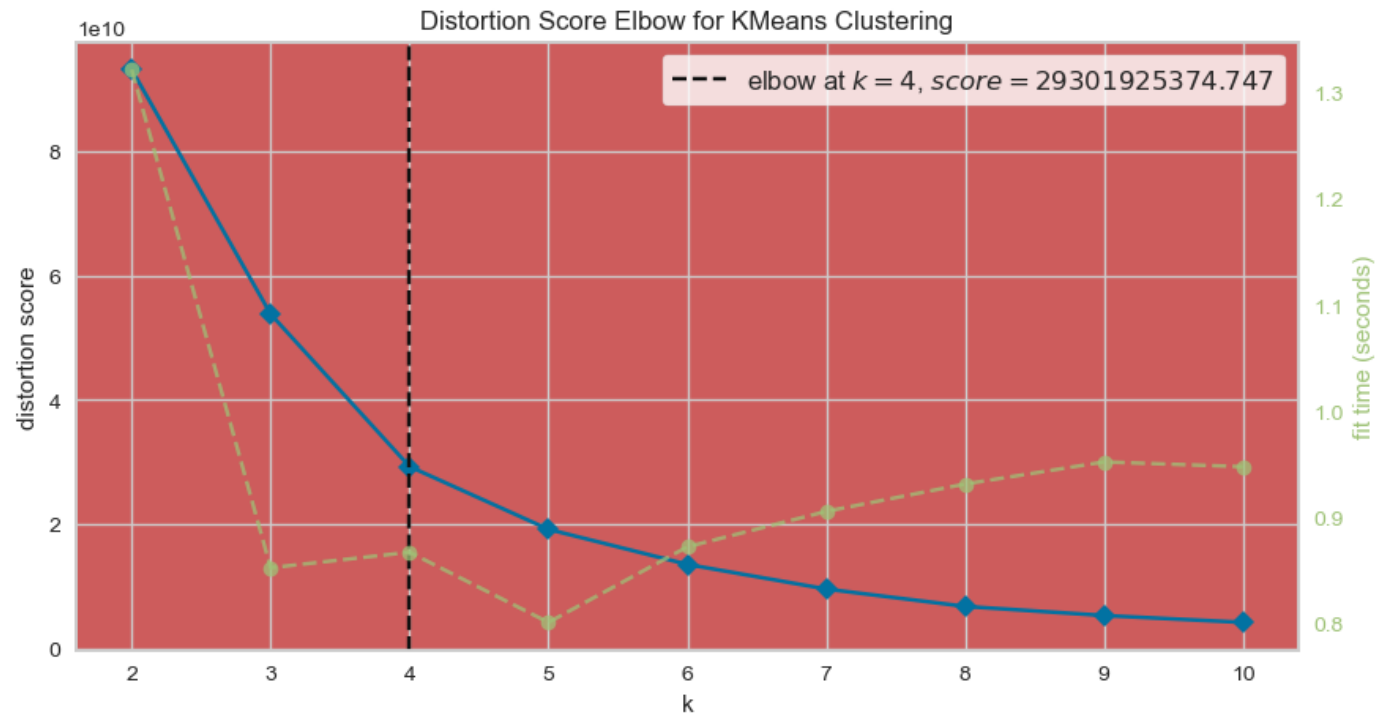
5. Modelling with KMeans Algorithm

```
In [25]: # Finding initial K value using Elbow Method
plt.figure(figsize=(10,5))
```

```

ax = plt.axes()
ax.set_facecolor("#cd5c5c")
Elbow_M = KElbowVisualizer(KMeans(), k=10)
Elbow_M.fit(RFMT)
Elbow_M.show()
print("Therefore K = 4")

```



Therefore K = 4

```

In [26]: #Fitting KMeans Model
kmeans = KMeans(n_clusters=4,max_iter=50)
kmeans.fit(RFMT)

```

```

Out[26]: KMeans
KMeans(max_iter=50, n_clusters=4)

```

```

In [27]: RFMT["Clusters"]=kmeans.labels_
RFMT.head().style.set_properties(**{"background-color": "#cd5c5c","color": "black", "bor

```

```

Out[27]: Recency Frequency Monetary InterPurchaseTime Clusters

```

CustomerID	Recency	Frequency	Monetary	InterPurchaseTime	Clusters
12346.000000	325	12	400.940000	33	0
12347.000000	2	8	4473.220000	50	0
12348.000000	75	5	779.730000	72	0
12349.000000	18	4	3347.990000	142	0
12352.000000	36	10	1739.490000	35	0

6. Model : Evaluation

```

In [28]: # how well the clusters are?:
#centriods
kmeans.cluster_centers_

```




```
Out[28]: array([[1.52262942e+02,  5.86186571e+00, 1.55692095e+03,  7.25317786e+01],
 [3.47727273e+01,  1.08227273e+02, 5.11814165e+04,  8.68181818e+00],
 [3.66666667e+00,  2.35333333e+02, 1.70986510e+05,  3.00000000e+00],
 [3.07256098e+01,  2.86250000e+01, 1.13062049e+04,  3.04542683e+01]])
```

```
In [29]: # grouping the data in accordance with each cluster seperately
one = RFMT[RFMT["Clusters"]==0]
two = RFMT[RFMT["Clusters"]==1]
three = RFMT[RFMT["Clusters"]==2]
four = RFMT[RFMT["Clusters"]==3]

#Checking the quality of clustering in the data set
plt.figure(figsize=(10,5))
ax = plt.axes()
ax.set_facecolor("#cd5c5c")
plt.scatter(one["Recency"],one["Frequency"],color='green')
plt.scatter(two["Recency"],two["Frequency"],color='red')
plt.scatter(three["Recency"],three["Frequency"],color='grey')
plt.scatter(four["Recency"],four["Frequency"],color='blue')
plt.scatter(kmeans.cluster_centers_[0,0],kmeans.cluster_centers_[0,1],color="yellow",marker='*')
plt.legend()
plt.show
```

```
Out[29]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
In [30]: from sklearn.metrics import silhouette_score
print("Silhouette score :",silhouette_score(RFMT, kmeans.labels_, metric='euclidean'))

Silhouette score : 0.797397771198647
```